




Figure 5 — Finite-sample inference: raw vs. bias-corrected leakage

A: raw $\hat{L}$ mean vs. n (dotted = population truth)B: bias-corrected $\tilde{L}$ mean ± 1 s.d. vs. n C: analytic bias $\beta_{pqr}(n)$ vs. empirical raw biasD: certified-boundary recovery probability vs. n 